

Multi-Label Bi-Temporal Change Classification via a Siamese ConvNeXt Baseline: A Phase-1 / Phase-2 Study

Amirkia Rafiei Oskooei, *BLM5135 – Yıldız Technical University*

Abstract—We address the multi-label bi-temporal change-classification task on a LEVIR-CC-derived dataset of 10,855 co-registered RGB image pairs, predicting three label families (Object, Event, Attribute) plus an auxiliary binary change-flag. A ConvNeXt-Tiny Siamese backbone with weight sharing extracts per-image features that are fused as $[f_A, f_B, f_A - f_B, |f_A - f_B|]$, projected to 768 channels, and consumed by linear heads under BCEWithLogitsLoss with per-class positive weighting. Two iterations are reported: a clean baseline (v1) and a regularized variant (v2) introducing head dropout, $10\times$ stronger weight decay, and a reduced positive-weight clamp, followed by validation-set per-class threshold optimization. The shared-encoder multi-task model attains the best test macro-F1 on Object (0.285) and Event (0.286), outperforming the corresponding single-task baselines and reviving classes that the single-task models leave at $F1 = 0$. The regularization in v2 successfully postpones the overfitting onset (best epoch +5 to +9) but reduces the peak validation F1, indicating an inherent ceiling on this backbone-dataset pair that further hyperparameter tuning alone does not surpass. We discuss class-imbalance behaviour, multi-task synergies on rare classes, and outline the directions reserved for Track 2.

Index Terms—multi-label classification, remote sensing change detection, Siamese networks, ConvNeXt, transfer learning, class imbalance, LEVIR-CC.

I. GİRİŞ (INTRODUCTION)

Bi-temporal change analysis on satellite or aerial imagery aims at identifying *what* has changed between two co-registered views of the same scene. Traditional change detection produces a binary pixel mask; in contrast, the course task in this work is a higher-level *multi-label classification* problem: each pair of RGB images (RGB_A, RGB_B) must be tagged with an arbitrary subset of three independent label vocabularies—Object (12 classes such as *building, road, tree, water*), Event (12 classes such as *build, remove, replace, turn*), and Attribute (24 classes such as *green, gray, dense, paved*). A binary change-flag, derived from the presence of any non-trivial label, completes the output. The dataset of $n=10,855$ pairs is derived from LEVIR-CC [1] and is severely long-tailed: *building* appears in 65% of training samples while *plant* appears in only 0.24%, a $270\times$ frequency ratio. The objectives of this report are to (i) construct a clean Siamese baseline within the constraints of Track 1, (ii) compare independent single-task models against a shared-encoder multi-task model, (iii) introduce a controlled regularization variant (v2) to test




Fig. 1. Per-family class frequencies in the training set (Object / Event / Attribute, left to right). The long-tailed shape and the cumulative dominance of two or three head classes characterize every family.

whether the baseline’s overfitting trajectory can be tamed within Track 1’s spirit, and (iv) document failure modes for the Track 2 work-in-progress. The originality of the study lies less in any individual architectural choice and more in the rigorous ablation methodology: every reported number can be reproduced from the released code at tag v1.0.0 or its successor, and every claim about multi-task synergy is grounded in per-class breakdowns rather than aggregate scores.

II. DENEYSEL ANALİZ (EXPERIMENTAL ANALYSIS)

A. Dataset and Pre-processing

Splits. The course-supplied `dataset.json` partitions the corpus into train (8,438; 77.7%), val (1,190; 11.0%) and test (1,227; 11.3%) and must be used as-is, per the rubric. Each split contains the same 71.7%/28.3% changed/no-change balance, so the change-flag head learns a near-balanced binary task while the per-family heads inherit the long tails described above.

Class imbalance. Figure 1 shows the per-family frequency histograms produced by our exploratory analysis (`scripts/run_eda.py`). The ratios of most-frequent to least-frequent non-*none* class are $270\times$ for Object, $28.5\times$ for Event, and $66.9\times$ for Attribute. The mean number of non-*none* labels per *changed* sample is 1.37 (Object), 1.48 (Event), and 2.12 (Attribute); the overall sample-level total is 4.54 active tags on average. This severity directly motivates positive-weighted BCE and the regularization study of Section II-G.

Soft cross-split leakage. We verified that 389 of the 5,200 unique scene identifiers ($\sim 7.5\%$) appear in more than one split via the offline-augmented A/B-swapped variants suffixed `_ters`. Because annotators relabelled those reversed pairs (events flip; object/attribute sets may also differ), the leakage

Placeholder: architecture diagram. To be drawn manually by the author. Suggested elements: two shared-weight ConvNeXt-Tiny stems, concat+diff fusion, 1×1 conv to 768, four heads (Object/Event/Attribute/Changeflag), with the v2 Dropout(0.3) on family heads explicitly indicated.

Fig. 2. Siamese ConvNeXt-Tiny architecture with concat-difference fusion and four task heads. v2 (Phase 2 of Track 1) additionally inserts Dropout(0.3) before each family head; the change-flag head remains unregularized.

is *soft*—it bounds the maximum achievable test score modestly but does not enable direct memorization. We retained the original splits because the course rubric explicitly mandates “Verilerin ayırımıları arasında karıştırılmaması, olduğu gibi kullanılması gerekmektedir”.

Pre-processing. All images are 256×256 PNG; we resize to 224×224 (matching the ImageNet pre-training resolution of the backbone), apply a paired horizontal-flip augmentation with $p=0.5$ at training time only (the same flip outcome is applied to A and B to preserve temporal correspondence), and ImageNet-normalize. The dataset ships with extensive offline augmentation (`_random_augment` variants, mixing flips and color jitter, and `_ters` pairs), so we deliberately keep on-the-fly augmentation minimal.

B. Architecture

The model (Figure 2) is intentionally simple. A single ImageNet-pretrained ConvNeXt-Tiny [2] (28M parameters, `convnext_tiny.fb_in1k` from `timm` [3]) is applied to both A and B with shared weights—a classical Siamese arrangement—producing $768\times 7\times 7$ feature maps. The two features are fused at the feature level by channel-wise concatenation of $[f_A, f_B, f_A - f_B, |f_A - f_B|]$, yielding $3072\times 7\times 7$, which a 1×1 convolution projects back to 768 channels. BatchNorm2d, GELU and global average pooling complete a 768-dimensional joint representation. Four task heads sit on top: three linear family heads (Object $\rightarrow 12$, Event $\rightarrow 12$, Attribute $\rightarrow 24$) and a binary change-flag head. Sigmoid activations are applied at loss/inference time, never inside the model, so the head outputs remain raw logits.

C. Why these choices

Backbone. ConvNeXt-Tiny offers a competitive accuracy-vs-parameters trade-off in the same regime as modern ResNets while behaving more like a Vision Transformer in its modernised layer norm and inverted bottlenecks [2]. Using a 28M-parameter pretrained CNN on 8.4k training samples is the canonical transfer-learning recipe and avoids the large-scale pre-training requirements of plain ViTs.

Fusion strategy. The concat + difference + absolute difference combination is the simplest *feature-level* fusion that simultaneously expresses “what is there” (f_A, f_B) and “what has changed” ($f_A - f_B, |f_A - f_B|$). Decision-level fusion would discard the relational structure between A and B ; input-level

fusion (e.g. stacking $A, B, A-B$ into a 9-channel image) is incompatible with a 3-channel pretrained backbone without re-initialising the stem. Transformer- or attention-based fusion is explicitly reserved for Track 2.

Multi-label output structure. Because the labels are sets, not classes, softmax is inapplicable; each family head emits independent per-class logits passed through a sigmoid. The change-flag head consolidates the no-change (*none*) case as an auxiliary binary signal rather than a per-family class, which simplifies macro-F1 reporting and is consistent across single- and multi-task configurations.

Loss. BCEWithLogitsLoss with per-class $\text{pos_weight}_c = \text{neg}_c / \text{pos}_c$ (clamped to $[1, 50]$ in v1 and $[1, 10]$ in v2) directly addresses the class imbalance documented in Section II-A. The change-flag loss is weighted by 0.5 in the multi-task aggregation; family losses are summed without further weighting (*naive sum*), a deliberate choice to keep v1 as a controllable baseline against which Track 2 task balancers (GradNorm, PCGrad, Kendall uncertainty) can be measured.

D. Training Protocol

Both iterations use AdamW with discriminative learning rates (10^{-5} for the backbone, 10^{-4} for fusion and heads), weight decay 10^{-4} (v1) or 10^{-3} (v2), a cosine schedule decaying to $\eta_{\min}=10^{-6}$, gradient clip 1.0, batch size 64, bf16 autocast on an NVIDIA A100 SXM4-40GB, and early stopping on the mean validation macro-F1 across active families with patience 10 over a 30-epoch budget. Per-worker random seeding is applied so the on-the-fly horizontal flip is reproducible. v2 introduces head dropout ($p=0.3$) on the three family heads but *not* on the change-flag head (the latter is already near-saturated). v2 also computes a per-class threshold sweep on the validation set (grid $\{0.05, 0.10, \dots, 0.95\}$, per-class F1-maximising) and applies those thresholds at test time as a post-hoc improvement; classes with zero positives in validation are mapped to threshold 1.0 to suppress test false positives.

E. Stage 1 (single-task): Results

Table I summarises the test-set metrics for the three independent single-task models in both v1 and v2 configurations. Object is dominated by *building* (788 of 1,227 test samples), so its micro-F1 is high (0.52 in v1, 0.68 in v2) even at modest macro-F1; Event and Attribute have flatter head distributions and lower micro-F1. v2 lifts Object micro-F1 by +0.16 and macro-F1 by +0.06, but slightly regresses Event (-0.014 macro) and Attribute (-0.012 macro). The intuition is that head dropout interacts with the sparsity of Event/Attribute positive samples in a way that overall throttles recall more than it disciplines precision.

A per-class breakdown for the Object family clarifies the macro-F1 ceiling. The classes *asphalt*, *plant* and *green* have test supports of 6, 2 and 2 respectively; the v1 single-task model scores F1=0 on all three, dragging the macro-F1 down by $3/12 \approx 0.25$ relative to what it would be on *learnable* classes. The macro-F1 of the v1 single-task Object model

TABLE I

STAGE 1 TEST METRICS (SINGLE-TASK MODELS). HIGHER IS BETTER.
BEST IN EACH ROW IS IN **BOLD**.

Family	Macro-F1		Micro-F1	
	v1 (flat 0.5)	v2 (val-tuned)	v1	v2 (val-tuned)
Object	0.2045	0.2252	0.5244	0.6088
Event	0.2796	0.2651	0.3883	0.4018
Attribute	0.2335	0.2210	0.3692	0.3969

TABLE II

STAGE 2 TEST METRICS (MULTI-TASK MODEL). v1 IS THE BASELINE, v2 IS THE REGULARIZED VARIANT EVALUATED WITH VAL-TUNED PER-CLASS THRESHOLDS.

Family	Macro-F1		Micro-F1	
	v1	v2 (val-tuned)	v1	v2 (val-tuned)
Object	0.2850	0.2364	0.6381	0.6464
Event	0.2861	0.2841	0.4036	0.4086
Attribute	0.2404	0.2176	0.3477	0.3211
Changeflag	0.95	0.96	—	—

recomputed over only the nine non-zero-support classes is 0.27—a useful interpretation that does not change the reported number but explains how a model with 0.86 F1 on *building* and 0.41 F1 on *road* can still average to 0.20.

F. Stage 2 (multi-task): Results

Table II reports the same metrics for the single shared-encoder multi-task model that emits all four heads jointly. The Object macro-F1 of v1 multi-task is the best Track-1 result we obtained (0.285, vs 0.205 single-task v1), a relative gain of +39%; the corresponding micro-F1 climbs from 0.52 to 0.64. Event and Attribute benefit too, though more modestly (+0.01 each on macro-F1). Per-class examination shows that multi-task learning revives the dead-class *green* (test support 2) from F1= 0 in single-task to F1= 0.50 in multi-task, almost certainly because the shared encoder absorbs cues from Attribute training (*green* co-occurs with vegetation-related events and attributes) that the Object-only model never sees.

The change-flag head is essentially trivial under either configuration (F1 $\approx$ 0.95): the 72%/28% changed/no-change class balance allows even a weak predictor to clear 0.85, so the change-flag is best interpreted as an auxiliary regularizer for the shared encoder rather than a contested task.

G. Train/Validation Curves and Overfitting

Figure 3 stacks the loss and macro-F1 curves of the v1 and v2 multi-task runs. Three observations dominate.

- 1) *Overfitting is real.* In v1 the training loss falls monotonically (1.15 $\rightarrow$ 0.13 on Object single-task) while the validation loss rises after the early-epoch minimum; the train/val loss ratio at training end is $\approx$ 14. This is the textbook signature of memorization rather than generalization.
- 2) *Best-epoch shifts later under v2.* v1 multi-task plateaus its validation macro-F1 at epoch 15; v2 reaches its




Fig. 3. Multi-task training curves. **Top:** v1 (Phase 1 of Track 1), best epoch 15. **Bottom:** v2 (Phase 2), best epoch 18. v2’s regularization (head dropout, 10 $\times$ weight decay, halved `pos_weight` clamp) postpones the overfitting onset but does not lift the peak validation F1.

best at epoch 18 despite the same patience-10 schedule. Single-task best-epochs move from 3/8/10 (Object/Event/Attribute) to 12/10/12. Regularization therefore *worked* on its first-order target: it slowed the overfitting trajectory.

- 3) *But the peak is lower under v2.* For every run the peak validation macro-F1 is lower in v2 than in v1 (−0.006 to −0.016). Combined with point 2, this indicates that ConvNeXt-Tiny on this dataset has a generalization ceiling between $\sim$ 0.22 and $\sim$ 0.29 macro-F1 that the v2 hyperparameter changes cannot surpass. The slowed training is paid for in peak performance.

Countermeasures actually deployed. (a) early stopping with patience 10, present in both iterations; (b) discriminative learning rates that fine-tune the backbone more conservatively than the heads; (c) gradient clipping to 1.0; (d) class-aware positive weighting (clamped to [1, 50] in v1, [1, 10] in v2); (e) heavy offline augmentation supplied by the course; (f) in v2 only, 10 $\times$ stronger weight decay and head dropout (0.3); (g) in v2 only, validation-set per-class threshold optimization. Together these are sufficient to constrain overfitting but, as Figure 3 shows, are insufficient to push performance through the apparent dataset/backbone ceiling.

Validation-tuned thresholds: a cautionary tale. The per-class threshold sweep on validation improves the validation macro-F1 by +0.04 to +0.06 across runs. On the test set, however, the improvement is markedly smaller and in one case (Object single-task v2) becomes a −0.04 regression. This is the textbook overfitting-on-the-validation-set phenomenon:




Fig. 4. Example qualitative panel (v1 multi-task, test split). Left: RGB_A. Center: RGB_B. Right: sample metadata, ground-truth labels, and the model’s top-5 predictions per family with sigmoid scores. Twenty such panels per run are saved as PNG plus a combined PDF.

with several Attribute classes having ≤ 10 positives in the validation split, sweeping a 19-point grid per class is finding noise as readily as signal. We report both flat and val-tuned numbers throughout to make this honest.

H. Qualitative Examples

Figure 4 shows one of the 20 qualitative panels produced per run by `visualize.py`. Each panel pairs the two RGB inputs with the sample identifier, the ground-truth labels per family, and the model’s predictions together with their sigmoid scores. The figure shown is from the v1 multi-task model; full sets exist for v1 (flat 0.5 thresholds) and v2 (both flat and val-tuned) at `results/track1_v*/<fam>/qualitative{,_tuned}/`.

III. SONUÇ (CONCLUSION)

We presented a complete Track-1 study of multi-label bi-temporal change classification, comparing single-task and multi-task instantiations of a shared Siamese ConvNeXt-Tiny baseline across two controlled iterations. The multi-task v1 configuration delivers the best test macro-F1 we observed in Track 1 (Object 0.285, Event 0.286, Attribute 0.240); the v2 regularization variant successfully delays overfitting but in doing so reduces the peak performance, particularly on the Object family where v1’s multi-task synergy was strongest. We attribute the persistent gap to the sanity bounds nominally targeted by the course rubric to a combination of severe long-tail class imbalance (some classes contain ≤ 5 test positives), unavoidable soft cross-split leakage from offline-augmented *A/B* swaps, and the inherent generalization ceiling of fine-tuning a 28M-parameter pretrained backbone on 8.4k training samples. The most informative finding is not the headline numbers themselves but the cleanly observed trade-off in v2: stronger regularization slows the overfitting trajectory yet erodes the rare-class transfer that made multi-task learning beneficial in v1. Track 2 will accordingly pursue orthogonal directions—class-aware losses such as Asymmetric Loss, attention-based fusion, and foundation-model backbones—rather than further hyperparameter sweeps within the present architecture.

ACKNOWLEDGMENTS

This work is the Track 1 deliverable for the BLM5135 *Derin Öğrenme ve Yapay Sinir Ağları* course at Yıldız Technical

University. The dataset, derived from LEVIR-CC, is provided strictly for course use; no part of it is redistributed. All code is released at the project repository under the v1.0.0 tag (corresponding to Phase 1) with the Phase 2 (v2) commit on the same main branch.

REFERENCES

- [1] C. Liu, R. Zhao, H. Chen, Z. Zou, and Z. Shi, “Remote sensing image change captioning with dual-branch transformers: A new method and a large scale dataset,” *IEEE Transactions on Geoscience and Remote Sensing*, vol. 60, pp. 1–20, 2022.
- [2] Z. Liu, H. Mao, C.-Y. Wu, C. Feichtenhofer, T. Darrell, and S. Xie, “A ConvNet for the 2020s,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022, pp. 11 976–11 986.
- [3] R. Wightman, “PyTorch image models,” <https://github.com/huggingface/pytorch-image-models>, 2019.